

Problem Set 6 (Total points: 110 + bonus 50)

Support Vector Machines

In this problem set you will implement an SVM and fit it using quadratic programming. We will use the CVXOPT module to solve the optimization problems.

You may want to start with solving the written problems at the end of this notebook or at least with reading the textbook. It will help a lot in this programming assignment.

Some of the cells will take minutes to run, so feel free to test your code on smaller tasks while you go. Easiest way would be to remove both for-loops and run the code just once.

Quadratic Programming

The standard form of a QP can be formulated as

$$\begin{aligned} \min_x \quad & \frac{1}{2}x^T Px + q^T x \\ \text{subject to} \quad & Gx \leq h \\ & Ax = b \end{aligned}$$

where \leq is an element-wise \leq . CVXOPT solver finds an optimal solution x^* , given a set of matrices P, q, G, h, A, b .

FYI, you can read on the methods to solve quadratic programming problems [here](https://en.wikipedia.org/wiki/Quadratic_programming#Solution_methods) (https://en.wikipedia.org/wiki/Quadratic_programming#Solution_methods).

Problem 1. [10 points]

Design appropriate matrices to solve the following problem.

$$\begin{aligned} \min_x \quad & f(x) = x_1^2 + 4x_2^2 - 8x_1 - 16x_2 \\ \text{subject to} \quad & x_1 + x_2 \leq 5 \\ & x_1 \leq 3 \\ & x_2 \geq 0 \end{aligned}$$

```
In [122]: from cvxopt import matrix, solvers
# Turns off the printing of CVXOPT solution for the rest of the notebook
solvers.options['show_progress'] = False

P = 2 * matrix([[1., 0.], [0., 4.]])
#-----

# Define q, G, h
# q =
# G =
# h =
#-----

q = -8*matrix([1.,2.])
G = matrix([[1.,1.,0.],[1.,0.,-1.]])
h = matrix([5.,3.,0.])
sol = solvers.qp(P, q, G, h)
x1, x2 = sol['x']
print('Optimal x: ({:.8f}, {:.8f})'.format(x1, x2))
```

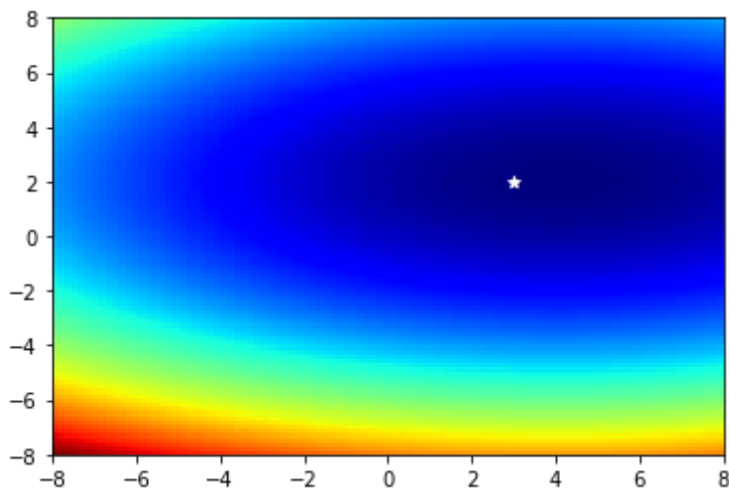
Optimal x: (2.99999993, 1.99927914)

Let's visualize the solution

```
In [123]: %matplotlib inline
from matplotlib import pyplot as plt
import numpy as np

X1, X2 = np.meshgrid(np.linspace(-8, 8, 100), np.linspace(-8, 8, 100))
F = X1**2 + 4*X2**2 - 8*X1 - 16 * X2

plt.pcolor(X1, X2, F, cmap='jet')
plt.scatter([x1], [x2], marker='*', color='white')
plt.show()
```



Why is the solution not in the minimum?

Answer:

The unconstrained local minimum lies outside of the allowable sub-space. Therefore the optimal solution cannot lie at the minimum as that would violate the constraint. Instead the solution lies on the constraint at the point nearest to the minimum.

Linear SVM

Now, let's implement linear SVM. We will do this for a general case, that allows class distributions to overlap (see Bishop 7.1.1).

As a linear model, linear SVM produces classification scores for a given sample x as

$$\hat{y}(x) = w^T \phi(x) + b$$

where $w \in \mathbb{R}^d$, $b \in \mathbb{R}$ are model weights and bias, respectively, and ϕ is a fixed feature-space transformation. Final label prediction is done by taking the sign of $\hat{y}(x)$.

Given a set of training samples $x_n \in \mathbb{R}^d$, $n = 1, \dots, N$, with the corresponding labels $y_i \in \{-1, 1\}$ linear SVM is fit (i.e. parameters w and b are chosen) by solving the following constrained optimization task:

$$\begin{aligned} \min_{w, \xi, b} \quad & \frac{1}{2} \|w\|^2 + C \sum_{n=1}^N \xi_n \\ \text{subject to} \quad & y_n \hat{y}(x_n) \geq 1 - \xi_n, \quad n = 1, \dots, N \\ & \xi_n \geq 0, \quad n = 1, \dots, N \end{aligned}$$

Problem 2.1 [60 points]

Your task is to implement this using a QP solver by designing appropriate matrices P , q , G , h .

Hints

1. You need to optimize over w, ξ, b . You can simply concatenate them into $\chi = (w, \xi, b)$ to feed it into QP-solver. Now, how to define the objective function and the constraints in terms of χ ? (For example, $b_1 + b_2$ can be obtained from vector $(a_1, b_1, b_2, c_1, c_2)$ by taking the inner product with $(0, 1, 1, 0, 0)$).
2. You can use `np.bmat` to construct matrices. Like this:

```
In [124]: np.bmat([[np.identity(3), np.zeros((3, 1))],  
                  [np.zeros((2, 3)), -np.ones((2, 1))]])
```

```
Out[124]: matrix([[ 1.,  0.,  0.,  0.],  
                  [ 0.,  1.,  0.,  0.],  
                  [ 0.,  0.,  1.,  0.],  
                  [ 0.,  0.,  0., -1.],  
                  [ 0.,  0.,  0., -1.]])
```

```

In [125]: from sklearn.base import BaseEstimator

class LinearSVM(BaseEstimator):
    def __init__(self, C, transform=None):
        self.C = C
        self.transform = transform

    def fit(self, X, Y):
        """Fit Linear SVM using training dataset (X, Y).

        :param X: data samples of shape (N, d).
        :param Y: data target labels of size (N). Each label is either 1 or -1.
        """

        # Apply transformation (phi) to X
        if self.transform is not None:
            X = self.transform(X)
        d = len(X[0])
        N = len(X)

        #-----

        # Construct appropriate matrices here to solve the optimization
        # problem described above.
        # We want optimal solution for vector (w, xi, b).

        P = np.bmat([[np.identity(d), np.zeros((d, N + 1))], [np.zeros((N
+ 1, N + d + 1))]])
        q = np.bmat([[np.zeros((d, 1))], [np.ones((N, 1))], [np.zeros((1, 1
))]])

        g1 = np.diag(Y) @ X @ np.bmat([np.identity(d), np.zeros((d, N+1
))])
        g2 = np.diag(Y) @ np.bmat([np.zeros((N, d + N)), np.ones((N, 1))])
        g3 = np.bmat([np.zeros((N, d)), np.identity(N), np.zeros((N, 1))])
        g = g1 + g2 + g3
        G = np.bmat([[g], [np.zeros((N, d)), np.identity(N), np.zeros((N, 1
))]])

        h = np.bmat([[np.ones((N, 1))], [np.zeros((N, 1))]])
        #-----

        P = matrix(P)
        q = matrix(self.C*q)
        G = matrix(-1*G)
        h = matrix(-1*h)

        sol = solvers.qp(P, q, G, h)
        ans = np.array(sol['x']).flatten()
        self.weights_ = ans[:d]
        self.xi_ = ans[d:d+N]
        self.bias_ = ans[-1]

        #-----

```

```

-----
        # Find support vectors. Must be a boolean array of length N havi
ng True for support
        # vectors and False for the rest.
#         margin = 1/np.sqrt(self.weights_.T@self.weights_)
        margin = Y*(X@self.weights_ + self.bias_)
        self.support_vectors = np.isclose(margin, 1 - self.xi_)
#-----

def predict_proba(self, X):
    """
    Make real-valued prediction for some new data.
    :param X: data samples of shape (N, d).
    :return: an array of N predicted scores.
    """
    if self.transform is not None:
        X = self.transform(X)

    y_hat = np.dot(X, self.weights_) + self.bias_
    return y_hat.flatten()

def predict(self, X):
    """
    Make binary prediction for some new data.
    :param X: data samples of shape (N, d).
    :return: an array of N binary predicted labels from {-1, 1}.
    """
    return np.sign(self.predict_proba(X))

```

Let's see how our LinearSVM performs on some data.

```

In [126]: from sklearn.datasets import make_classification, make_circles
X = [None, None, None]
y = [None, None, None]
X[0], y[0] = make_classification(n_samples=100, n_features=2, n_redundan
t=0, n_clusters_per_class=1, random_state=1)
X[1], y[1] = make_circles(n_samples=100, factor=0.5)
X[2], y[2] = make_classification(n_samples=100, n_features=2, n_redundan
t=0, n_clusters_per_class=1, random_state=4)

# Go from {0, 1} to {-1, 1}
y = [2 * yy - 1 for yy in y]

```

```

In [127]: C_values = [0.01, 0.1, 1.]

plot_i = 0
plt.figure(figsize=(len(X) * 7, len(C_values) * 7))
for C in C_values:
    for i in range(len(X)):
        plot_i += 1
        plt.subplot(len(C_values), len(X), plot_i)
        #-----

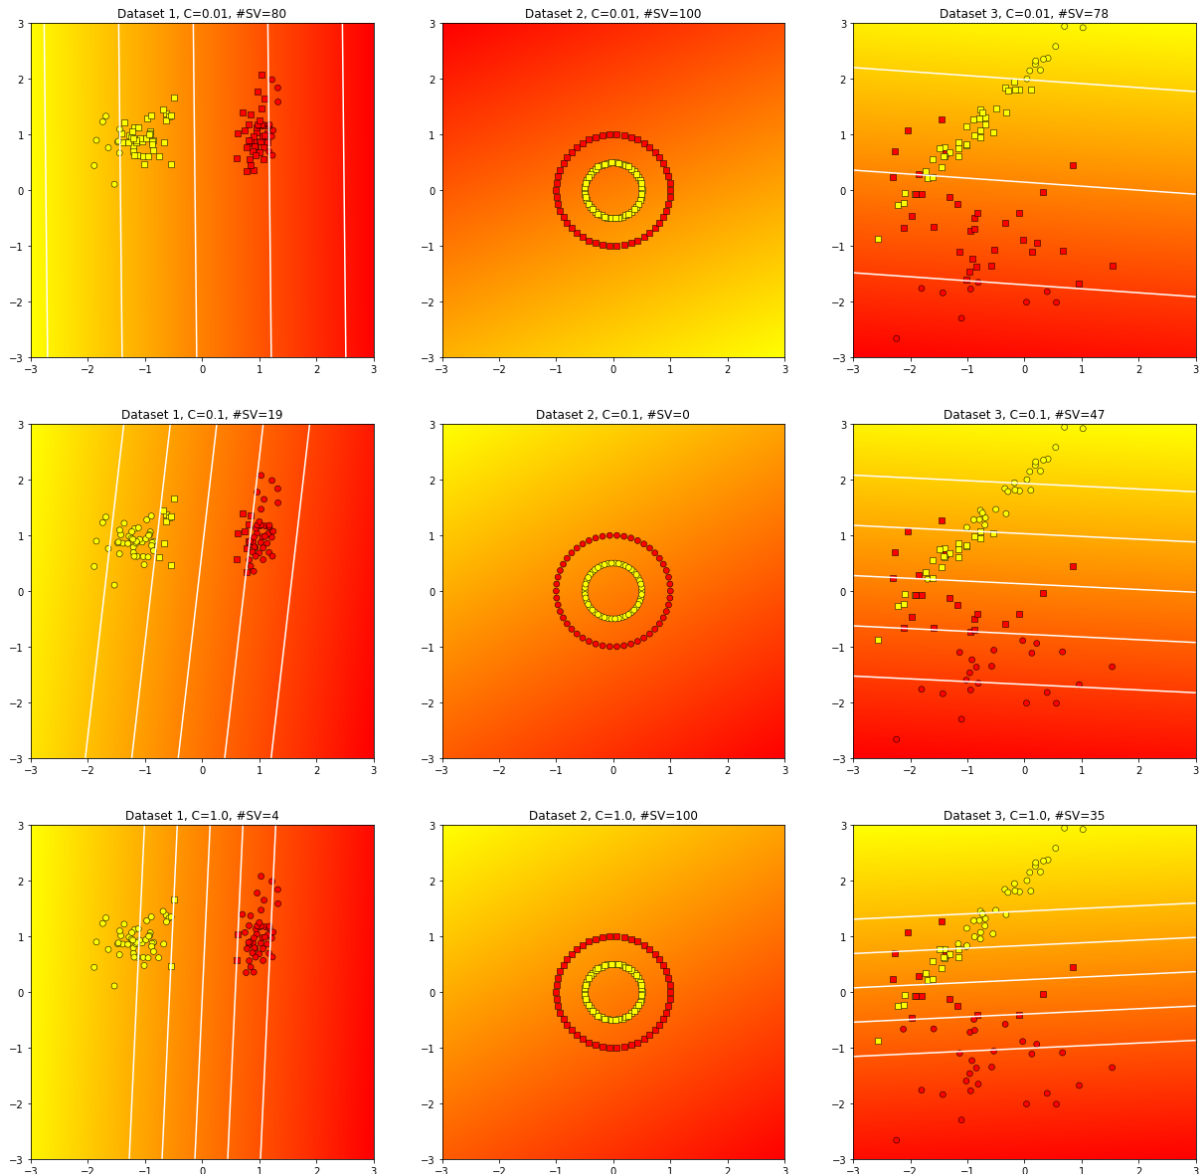
        model = LinearSVM(C=C)
        #-----

        model.fit(X[i], y[i])
        sv = model.support_vectors
        n_sv = sv.sum()
        if n_sv > 0:
            plt.scatter(X[i][:, 0][sv], X[i][:, 1][sv], c=y[i][sv], cmap
='autumn', marker='s',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        if n_sv < len(X[i]):
            plt.scatter(X[i][:, 0][~sv], X[i][:, 1][~sv], c=y[i][~sv], c
map='autumn',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        xvals = np.linspace(-3, 3, 200)
        yvals = np.linspace(-3, 3, 200)
        xx, yy = np.meshgrid(xvals, yvals)
        zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel
()]), xx.shape)
        plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
        plt.contour(xx, yy, zz, levels=(-2,-1, 0, 1,2), colors='w', line
widths=1.5, zorder=1, linestyle='solid')

        plt.xlim([-3, 3])
        plt.ylim([-3, 3])
        plt.title('Dataset {}, C={}, #SV={}'.format(i + 1, C, n_sv))
plt.show()

```

```
/Users/danielvaroli/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:26: UserWarning: No contour levels were found within the data range.
```



Why does the number of support vectors decrease as C increases?

Answer:

As we increase the value of C , we increase the penalty of mis-classification, as a result we bring the margin closer to the decision boundary as we are being more restrictive. Hence as the margin gets smaller the number of support vectors decreases.

For debug purposes. Very last model must have almost the same weights and bias:

$$w = \begin{pmatrix} -0.0784521 \\ 1.62264867 \end{pmatrix}$$

$$b = -0.3528510092782581$$


```
In [128]: model.weights_
```

```
Out[128]: array([-0.0784521 ,  1.62264867])
```

```
In [129]: model.bias_
```

```
Out[129]: -0.3528510092782581
```

Problem 2.2 [10 points]

Even using a linear SVM, we are able to separate data that is linearly inseparable by using feature transformation.

Implement the following feature transformation $\phi(x_1, x_2) = (x_1, x_2, x_1^2, x_2^2, x_1x_2)$

```
In [130]: def append_second_order(X):  
    """Given array Nx[x1, x2] return Nx[x1, x2, x1^2, x2^2, x1x2]."""  
    x1 = np.array(X[:,0])  
    x2 = np.array(X[:,1])  
    new_X = np.column_stack((x1,x2,x1**2,x2**2,x1*x2))  
    return new_X  
  
    assert np.all(append_second_order(np.array([[1, 2]])) == np.array([[1, 2  
    , 1, 4, 2]])), 'Transformation is incorrect.'
```

```

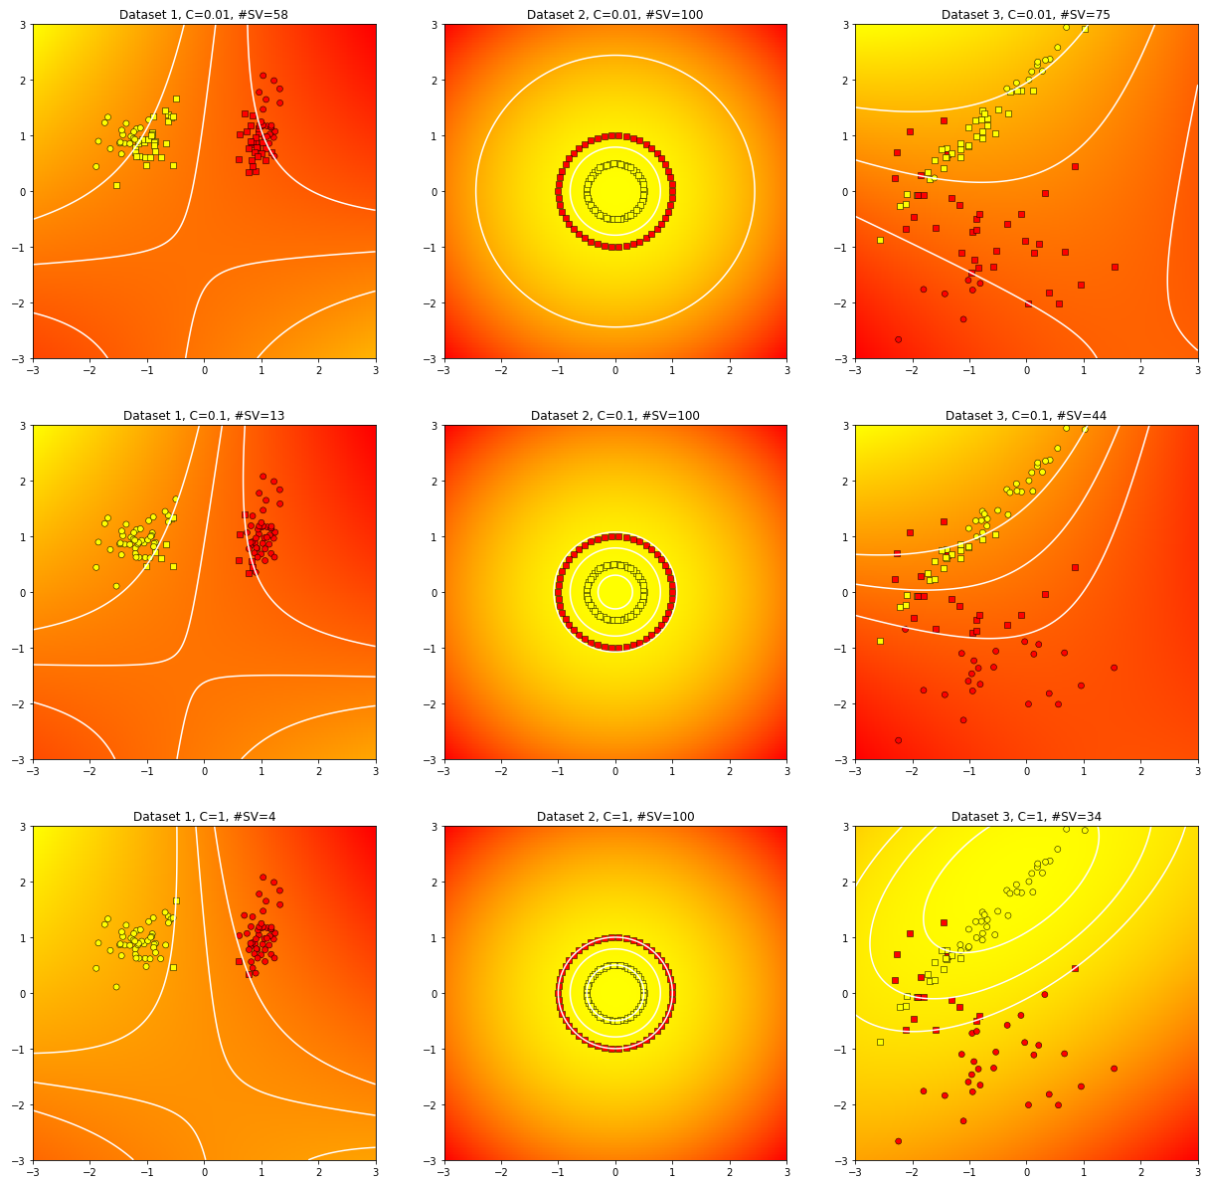
In [131]: plot_i = 0
C_values = [0.01, 0.1, 1]
plt.figure(figsize=(len(X) * 7, len(C_values) * 7))
for C in C_values:
    for i in range(len(X)):
        plot_i += 1
        plt.subplot(len(C_values), len(X), plot_i)
        #-----

        model = LinearSVM(C=C, transform=append_second_order)
        #-----

        model.fit(X[i], y[i])
        sv = model.support_vectors
        n_sv = sv.sum()
        if n_sv > 0:
            plt.scatter(X[i][:, 0][sv], X[i][:, 1][sv], c=y[i][sv], cmap
='autumn', marker='s',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        if n_sv < len(X[i]):
            plt.scatter(X[i][:, 0][~sv], X[i][:, 1][~sv], c=y[i][~sv], c
map='autumn',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        xvals = np.linspace(-3, 3, 200)
        yvals = np.linspace(-3, 3, 200)
        xx, yy = np.meshgrid(xvals, yvals)
        zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel
()]), xx.shape)
        plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
        plt.contour(xx, yy, zz, levels=(-1, 0, 1), colors='w', linewidth
hs=1.5, zorder=1, linestyle='solid')

        plt.xlim([-3, 3])
        plt.ylim([-3, 3])
        plt.title('Dataset {}, C={}, #SV={}'.format(i + 1, C, n_sv))
plt.show()

```



Bonus part (Optional)

Dual representation. Kernel SVM

The dual representation of the maximum margin problem is given by

$$\begin{aligned} \max_{\alpha} \quad & \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m y_n y_m k(x_n, x_m) \\ \text{subject to} \quad & 0 \leq \alpha_n \leq C, \quad n = 1, \dots, N \\ & \sum_{n=1}^N \alpha_n y_n = 0 \end{aligned}$$

In this case bias b can be computed as

$$b = \frac{1}{|S|} \sum_{n \in S} \left(y_n - \sum_{m \in S} \alpha_m y_m k(x_n, x_m) \right),$$

and the prediction turns into

$$\hat{y}_i(x) = \sum_{n \in S} \alpha_n y_n k(x_n, x_i) + b.$$

Everywhere above k is a kernel function: $k(x_1, x_2) = \phi(x_1)^T \phi(x_2)$ (and the trick is that we don't have to specify ϕ , just k).

Note, that now

1. We want to maximize the objective function, not minimize it.
2. We have equality constraints. (That means we should use A and b in qp-solver)
3. We need access to the support vectors (but not all the training samples) in order to make a prediction.

Problem 3.1 [40 points]

Implement KernelSVM

Hints

1. What is the variable we are optimizing over? lagrange multipliers
2. How can we maximize a function given a tool for minimization? Take the negative
3. What is the definition of a support vector in the dual representation?

```

In [193]: class KernelSVM(BaseEstimator):
    def __init__(self, C, kernel=linear_kernel):
        self.C = C
        self.kernel = kernel

    def fit(self, X, Y):
        """Fit Kernel SVM using training dataset (X, Y).

        :param X: data samples of shape (N, d).
        :param Y: data target labels of size (N). Each label is either 1
        or -1. Denoted as  $t_i$  in Bishop.
        """
        N, num_features = X.shape

        K = np.zeros((N, N))
        for i in range(N):
            for j in range(N):
                K[i,j] = self.kernel(X[i], X[j])

        #-----
        # Construct appropriate matrices here to solve the optimization
        # problem described above.
        P = matrix(np.outer(Y,Y) * K)
        q = matrix(-1. * np.ones(N) )
        A = 1. * matrix(Y.reshape((1,-1)))
        b = matrix(0.0)
        G = matrix(np.bmat([[ -1.*np.identity(N)], [ 1.*np.identity(N) ]]))
        h = matrix(np.bmat([[ np.zeros((N,1)) ], [ 1.*self.C*np.ones((N,1
        )) ]]))

        #-----

        solution = solvers.qp(P, q, G, h, A, b)
        self.alpha_ = np.ravel(solution['x'])

        #-----
        # Find support vectors. Must be a boolean array of length N havi
        # ng True for support
        # vectors and False for the rest.
        self.support_vectors = self.alpha_ > 1e-5
        #-----

        #         ind = np.arange(len(a))[sv]

        sv_ind = self.support_vectors.nonzero()[0]
        self.alpha_sup = self.alpha_[sv_ind]
        self.X_sup = X[sv_ind]
        self.Y_sup = Y[sv_ind]
        self.n_sv = len(sv_ind)
        print("Num Support Vectors: %d" % self.n_sv)

```

```

#-----
# Compute bias

self.bias_ = 0
for n in range(len(self.alpha_sup)):
    self.bias_ += self.Y_sup[n]
    self.bias_ -= np.sum(self.alpha_sup * self.Y_sup * K[sv_ind[
n],self.support_vectors])

self.bias_ /= self.n_sv

#-----

def predict_proba(self, X):
    """
    Make real-valued prediction for some new data.
    :param X: data samples of shape (N, d).
    :return: an array of N predicted scores.
    """
    y_predict = np.zeros(len(X))
    for i in range(len(X)):
        s = 0
        for alpha_sup, Y_sup, X_sup in zip(self.alpha_sup, self.Y_sup, self.X_sup):
            s += alpha_sup * Y_sup * self.kernel(X[i], X_sup)
        y_predict[i] = s

    return y_predict + self.bias_

def predict(self, X):
    """
    Make binary prediction for some new data.
    :param X: data samples of shape (N, d).
    :return: an array of N binary predicted labels from {-1, 1}.
    """
    return np.sign(self.predict_proba(X))

```

We can first test our implementation by using the dot product as a kernel function. What should we expect in this case?

```

In [194]: def linear_kernel(x, y):
            return np.dot(x, y)

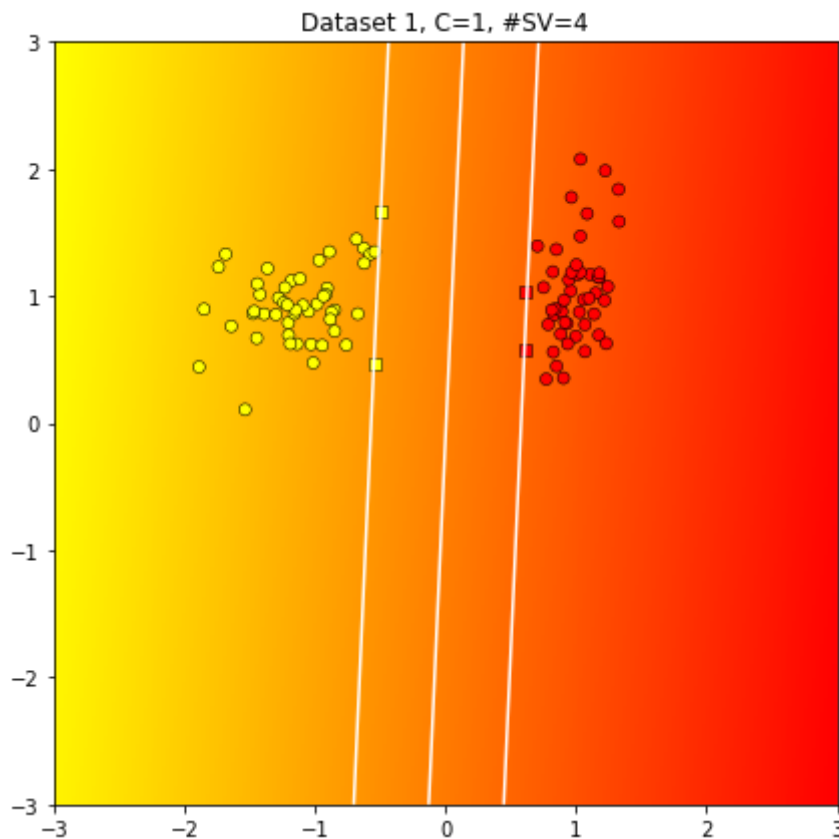
C = 1
i = 0

plt.figure(figsize=(7, 7))
#-----
-----
model = KernelSVM(C=C, kernel=linear_kernel)
#-----
-----
model.fit(X[i], y[i])
sv = model.support_vectors ##### EDITED ##### should be model.support_vectors
n_sv = sv.sum()
if n_sv > 0:
    plt.scatter(X[i][:, 0][sv], X[i][:, 1][sv], c=y[i][sv], cmap='autumn', marker='s',
                linewidths=0.5, edgecolors=(0, 0, 0, 1))
if n_sv < len(X[i]):
    plt.scatter(X[i][:, 0][~sv], X[i][:, 1][~sv], c=y[i][~sv], cmap='autumn',
                linewidths=0.5, edgecolors=(0, 0, 0, 1))
xvals = np.linspace(-3, 3, 200)
yvals = np.linspace(-3, 3, 200)
xx, yy = np.meshgrid(xvals, yvals)
zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel()])), xx.shape)
plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
plt.contour(xx, yy, zz, levels=(-1, 0, 1), colors='w', linewidths=1.5,
            zorder=1, linestyle='solid')

plt.xlim([-3, 3])
plt.ylim([-3, 3])
plt.title('Dataset {}, C={}, #SV={}'.format(i + 1, C, n_sv))
plt.show()

```

Num Support Vectors: 4



Problem 3.2 [5 points]

Implement a polynomial kernel function ([wiki \(https://en.wikipedia.org/wiki/Polynomial_kernel\)](https://en.wikipedia.org/wiki/Polynomial_kernel)).

```
In [195]: def polynomial_kernel(d, c=0):
            """Returns a polynomial kernel FUNCTION."""
            def kernel(x, y):
                """
                :param x: vector of size L
                :param y: vector of size L
                :return: [polynomial kernel of degree d with bias parameter c] o
                f x and y. A scalar.
                """
                return (np.dot(x.T, y) + c)**d
            return kernel

            assert polynomial_kernel(d=2, c=1)(np.array([1, 2]), np.array([3, 4])) =
            = 144, 'Polynomial kernel implemented incorrectly'
```

Let's see how it performs. This might take some time to run.


```

In [196]: plot_i = 0
C = 10
d_values = [2,3,4]
plt.figure(figsize=(len(X) * 7, len(d_values) * 7))
for d in d_values:
    for i in range(len(X)):
        plot_i += 1
        plt.subplot(len(d_values), len(X), plot_i)
        #-----

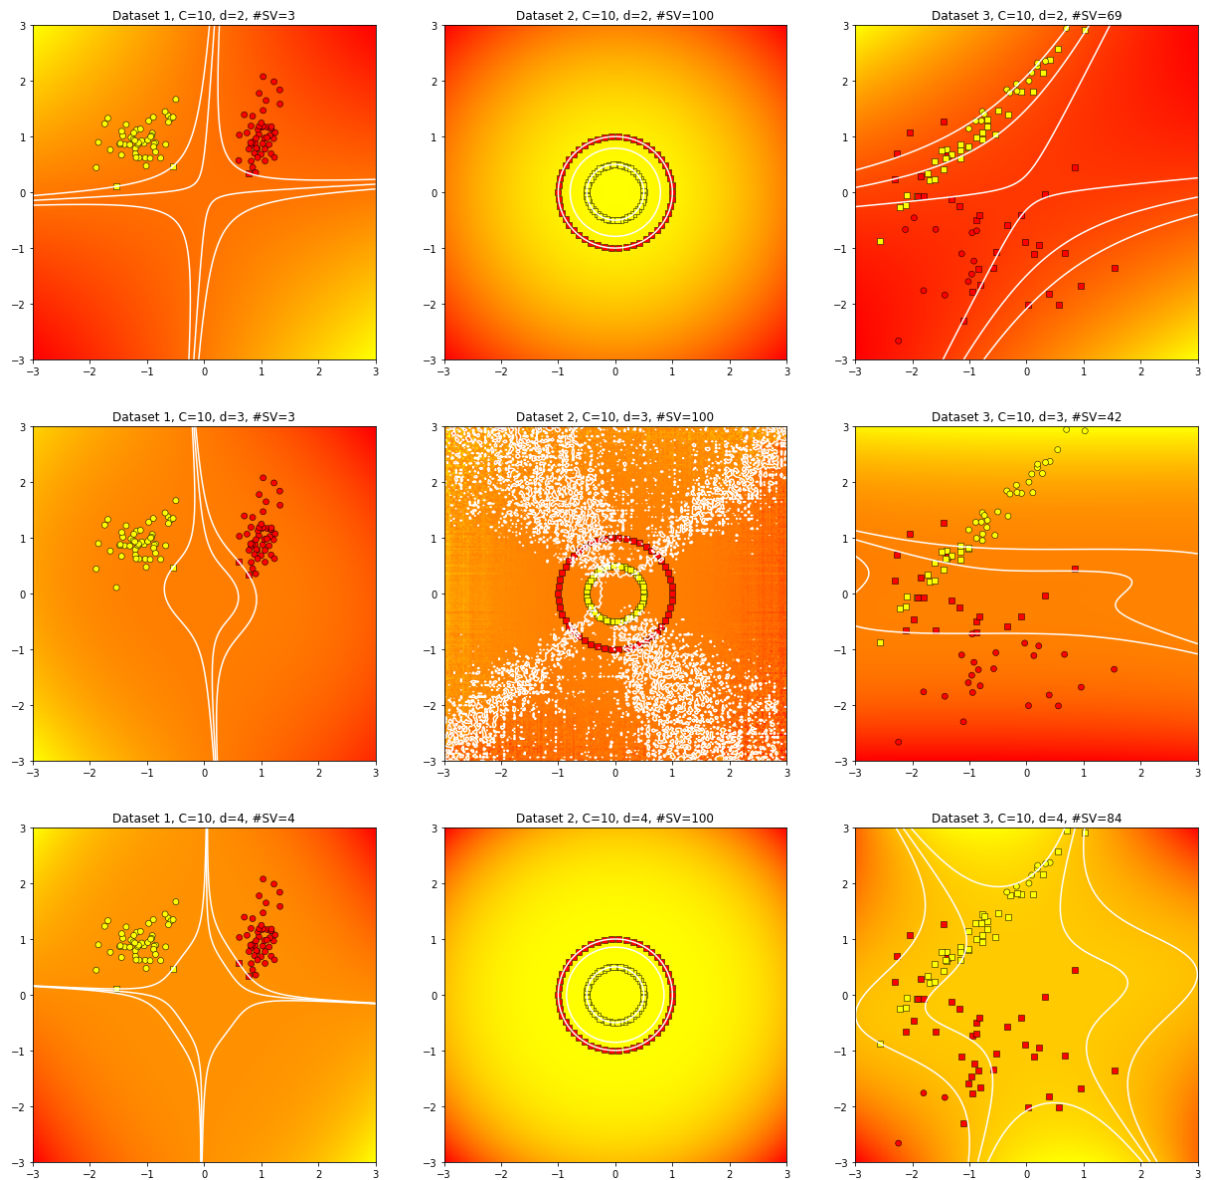
        model = KernelSVM(C=C, kernel=polynomial_kernel(d))
        #-----

        model.fit(X[i], y[i])
        sv = model.support_vectors
        n_sv = sv.sum()
        if n_sv > 0:
            plt.scatter(X[i][:, 0][sv], X[i][:, 1][sv], c=y[i][sv], cmap
='autumn', marker='s',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        if n_sv < len(X[i]):
            plt.scatter(X[i][:, 0][~sv], X[i][:, 1][~sv], c=y[i][~sv], c
map='autumn',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        xvals = np.linspace(-3, 3, 200)
        yvals = np.linspace(-3, 3, 200)
        xx, yy = np.meshgrid(xvals, yvals)
        zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel
()]), xx.shape)
        plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
        plt.contour(xx, yy, zz, levels=(-1, 0, 1), colors='w', linewidth
hs=1.5, zorder=1, linestyle='solid')

        plt.xlim([-3, 3])
        plt.ylim([-3, 3])
        plt.title('Dataset {}, C={}, d={}, #SV={}'.format(i + 1, C, d, n
_sv))

```

Num Support Vectors: 3
 Num Support Vectors: 100
 Num Support Vectors: 69
 Num Support Vectors: 3
 Num Support Vectors: 100
 Num Support Vectors: 42
 Num Support Vectors: 4
 Num Support Vectors: 100
 Num Support Vectors: 84



Task 3.3 [5 points]

Finally, you need to implement a **radial basis function** kernel ([wiki](https://en.wikipedia.org/wiki/Radial_basis_function_kernel) (https://en.wikipedia.org/wiki/Radial_basis_function_kernel)).

```
In [197]: def RBF_kernel(sigma):  
    """Returns an RBF kernel FUNCTION."""  
    def kernel(x, y):  
        """  
        :param x: vector of size L  
        :param y: vector of size L  
        :return: [rbf kernel with parameter sigma] of x and y. A scalar.  
        """  
        return np.exp(-np.linalg.norm(x-y)**2 / (2 * (sigma ** 2)))  
    return kernel
```

Let's see how it performs. This might take some time to run.

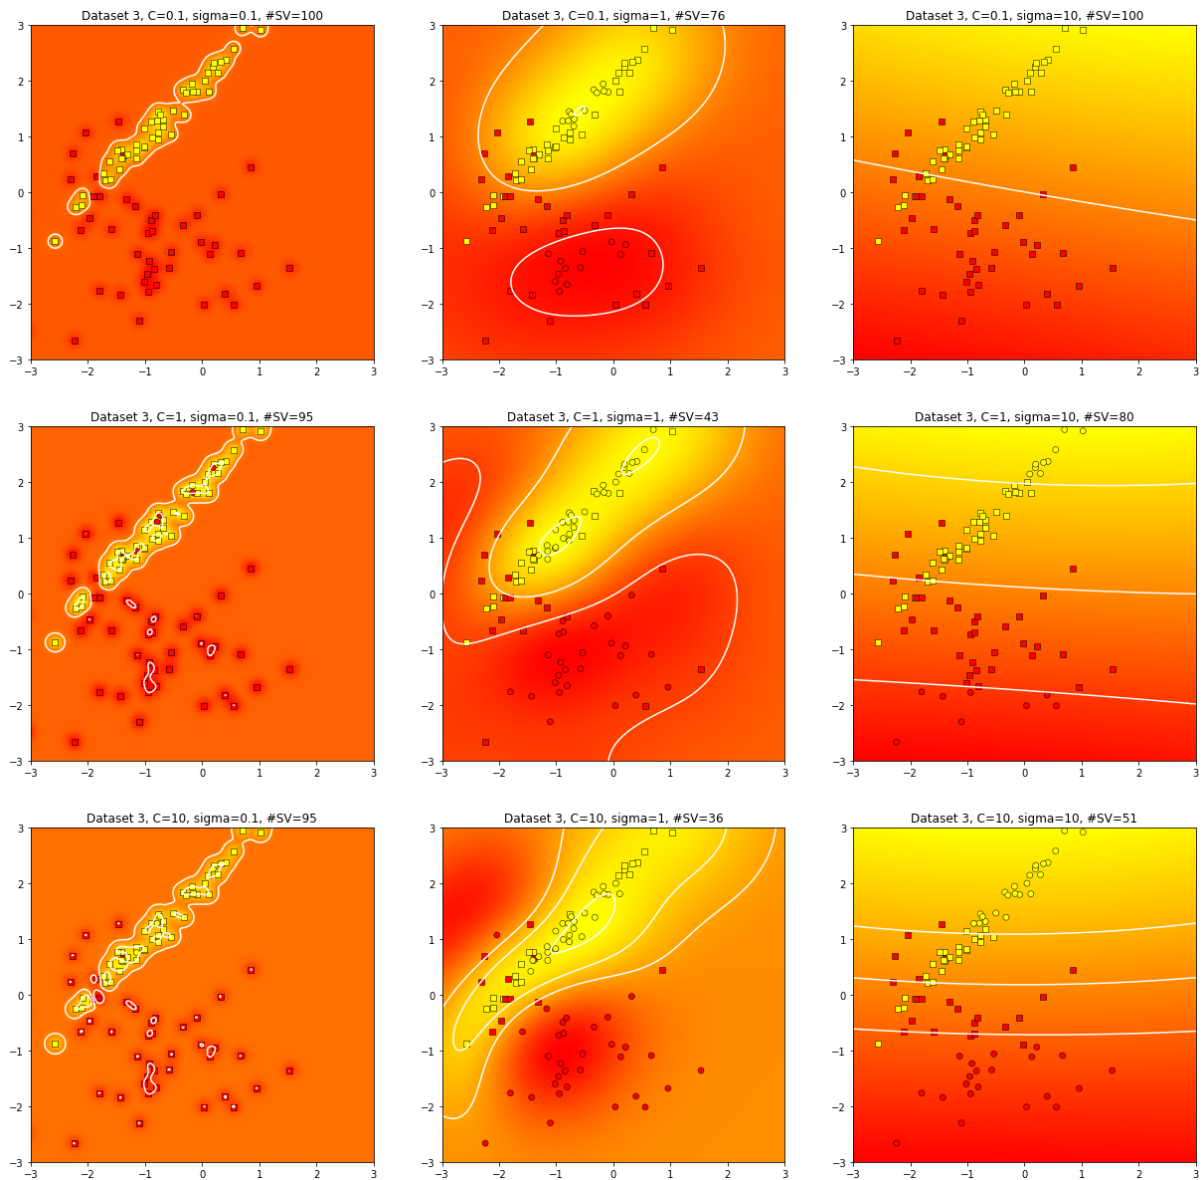
```

In [188]: plot_i = 0
C_values = [0.1, 1, 10]
sigma_values = [0.1, 1, 10]
plt.figure(figsize=(len(sigma_values) * 7, len(C_values) * 7))
i = 2
for C in C_values:
    for sigma in sigma_values:
        plot_i += 1
        plt.subplot(len(C_values), len(X), plot_i)
        model = KernelSVM(C=C, kernel=RBF_kernel(sigma))
        model.fit(X[i], y[i])
        sv = model.support_vectors
        n_sv = sv.sum()
        if n_sv > 0:
            plt.scatter(X[i][:, 0][sv], X[i][:, 1][sv], c=y[i][sv], cmap
='autumn', marker='s',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        if n_sv < len(X[i]):
            plt.scatter(X[i][:, 0][~sv], X[i][:, 1][~sv], c=y[i][~sv], c
map='autumn',
                        linewidths=0.5, edgecolors=(0, 0, 0, 1))
        xvals = np.linspace(-3, 3, 200)
        yvals = np.linspace(-3, 3, 200)
        xx, yy = np.meshgrid(xvals, yvals)
        zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel
()]), xx.shape)
        plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
        plt.contour(xx, yy, zz, levels=(-1, 0, 1, ), colors='w', linewidth
hs=1.5, zorder=1, linestyle='solid')

        plt.xlim([-3, 3])
        plt.ylim([-3, 3])
        plt.title('Dataset {}, C={}, sigma={}, #SV={}'.format(i + 1, C,
sigma, n_sv))

```

100 support vectors out of 100 points
 76 support vectors out of 100 points
 100 support vectors out of 100 points
 95 support vectors out of 100 points
 43 support vectors out of 100 points
 80 support vectors out of 100 points
 95 support vectors out of 100 points
 36 support vectors out of 100 points
 51 support vectors out of 100 points



Well done!

Awesome! Now you understand all of the important parameters in SVMs. Have a look at SVM from scikit-learn module and how it is used (very similar to ours).

```
In [17]: from sklearn.svm import SVC  
SVC?
```

Init signature: SVC(C=1.0, kernel='rbf', degree=3, gamma='auto_deprecated', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

Docstring:

C-Support Vector Classification.

The implementation is based on libsvm. The fit time complexity is more than quadratic with the number of samples which makes it hard to scale to dataset with more than a couple of 10000 samples.

The multiclass support is handled according to a one-vs-one scheme.

For details on the precise mathematical formulation of the provided kernel functions and how `gamma`, `coef0` and `degree` affect each other, see the corresponding section in the narrative documentation: :ref:`svm_kernels`.

Read more in the :ref:`User Guide <svm_classification>`.

Parameters

C : float, optional (default=1.0)
Penalty parameter C of the error term.

kernel : string, optional (default='rbf')
Specifies the kernel type to be used in the algorithm.
It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or
a callable.
If none is given, 'rbf' will be used. If a callable is given it is
used to pre-compute the kernel matrix from data matrices; that matrix
should be an array of shape ``(n_samples, n_samples)``.

degree : int, optional (default=3)
Degree of the polynomial kernel function ('poly').
Ignored by all other kernels.

gamma : float, optional (default='auto')
Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.

Current default is 'auto' which uses $1 / n_{\text{features}}$,
if ``gamma='scale'`` is passed then it uses $1 / (n_{\text{features}} * X.\text{std}())$
as value of gamma. The current default of gamma, 'auto', will change
to 'scale' in version 0.22. 'auto_deprecated', a deprecated version
of
'auto' is used as a default indicating that no explicit value of gamma
was passed.

coef0 : float, optional (default=0.0)
Independent term in kernel function.
It is only significant in 'poly' and 'sigmoid'.

`shrinking` : boolean, optional (default=True)
Whether to use the shrinking heuristic.

`probability` : boolean, optional (default=False)
Whether to enable probability estimates. This must be enabled prior to calling ``fit``, and will slow down that method.

`tol` : float, optional (default=1e-3)
Tolerance for stopping criterion.

`cache_size` : float, optional
Specify the size of the kernel cache (in MB).

`class_weight` : {dict, 'balanced'}, optional
Set the parameter C of class i to `class_weight[i]*C` for SVC. If not given, all classes are supposed to have weight one.
The "balanced" mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as ``n_samples / (n_classes * np.bincount(y))``

`verbose` : bool, default: False
Enable verbose output. Note that this setting takes advantage of a per-process runtime setting in libsvm that, if enabled, may not work properly in a multithreaded context.

`max_iter` : int, optional (default=-1)
Hard limit on iterations within solver, or -1 for no limit.

`decision_function_shape` : 'ovo', 'ovr', default='ovr'
Whether to return a one-vs-rest ('ovr') decision function of shape (n_samples, n_classes) as all other classifiers, or the original one-vs-one ('ovo') decision function of libsvm which has shape (n_samples, n_classes * (n_classes - 1) / 2). However, one-vs-one ('ovo') is always used as multi-class strategy.

.. versionchanged:: 0.19
 `decision_function_shape` is 'ovr' by default.

.. versionadded:: 0.17
 `*decision_function_shape='ovr'` is recommended.

.. versionchanged:: 0.17
 Deprecated `*decision_function_shape='ovo'` and `None`.

`random_state` : int, RandomState instance or None, optional (default=None)
The seed of the pseudo random number generator used when shuffling the data for probability estimates. If int, `random_state` is the seed used by the random number generator; If RandomState instance, `random_state` is the random number generator; If None, the random number generator is the RandomState instance used by ``np.random``.

Attributes

```

support_ : array-like, shape = [n_SV]
    Indices of support vectors.

support_vectors_ : array-like, shape = [n_SV, n_features]
    Support vectors.

n_support_ : array-like, dtype=int32, shape = [n_class]
    Number of support vectors for each class.

dual_coef_ : array, shape = [n_class-1, n_SV]
    Coefficients of the support vector in the decision function.
    For multiclass, coefficient for all 1-vs-1 classifiers.
    The layout of the coefficients in the multiclass case is somewhat
    non-trivial. See the section about multi-class classification in the
    SVM section of the User Guide for details.

coef_ : array, shape = [n_class * (n_class-1) / 2, n_features]
    Weights assigned to the features (coefficients in the primal
    problem). This is only available in the case of a linear kernel.

    `coef_` is a readonly property derived from `dual_coef_` and
    `support_vectors_`.

intercept_ : array, shape = [n_class * (n_class-1) / 2]
    Constants in decision function.

```

Examples

```

-----
>>> import numpy as np
>>> X = np.array([[ -1, -1], [-2, -1], [1, 1], [2, 1]])
>>> y = np.array([1, 1, 2, 2])
>>> from sklearn.svm import SVC
>>> clf = SVC(gamma='auto')
>>> clf.fit(X, y) #doctest: +NORMALIZE_WHITESPACE
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
>>> print(clf.predict([[ -0.8, -1]]))
[1]

```

See also

SVR

Support Vector Machine for Regression implemented using libsvm.

LinearSVC

Scalable Linear Support Vector Machine for classification implemented using liblinear. Check the See also section of LinearSVC for more comparison element.

File: c:\miniconda\envs\cvx\lib\site-packages\sklearn\svm\classes.py
Type: ABCMeta

```

In [18]: plot_i = 0
C = 10
d_values = [2, 3, 4]
plt.figure(figsize=(len(X) * 7, len(d_values) * 7))
for d in d_values:
    for i in range(len(X)):
        plot_i += 1
        plt.subplot(len(d_values), len(X), plot_i)
        #-----

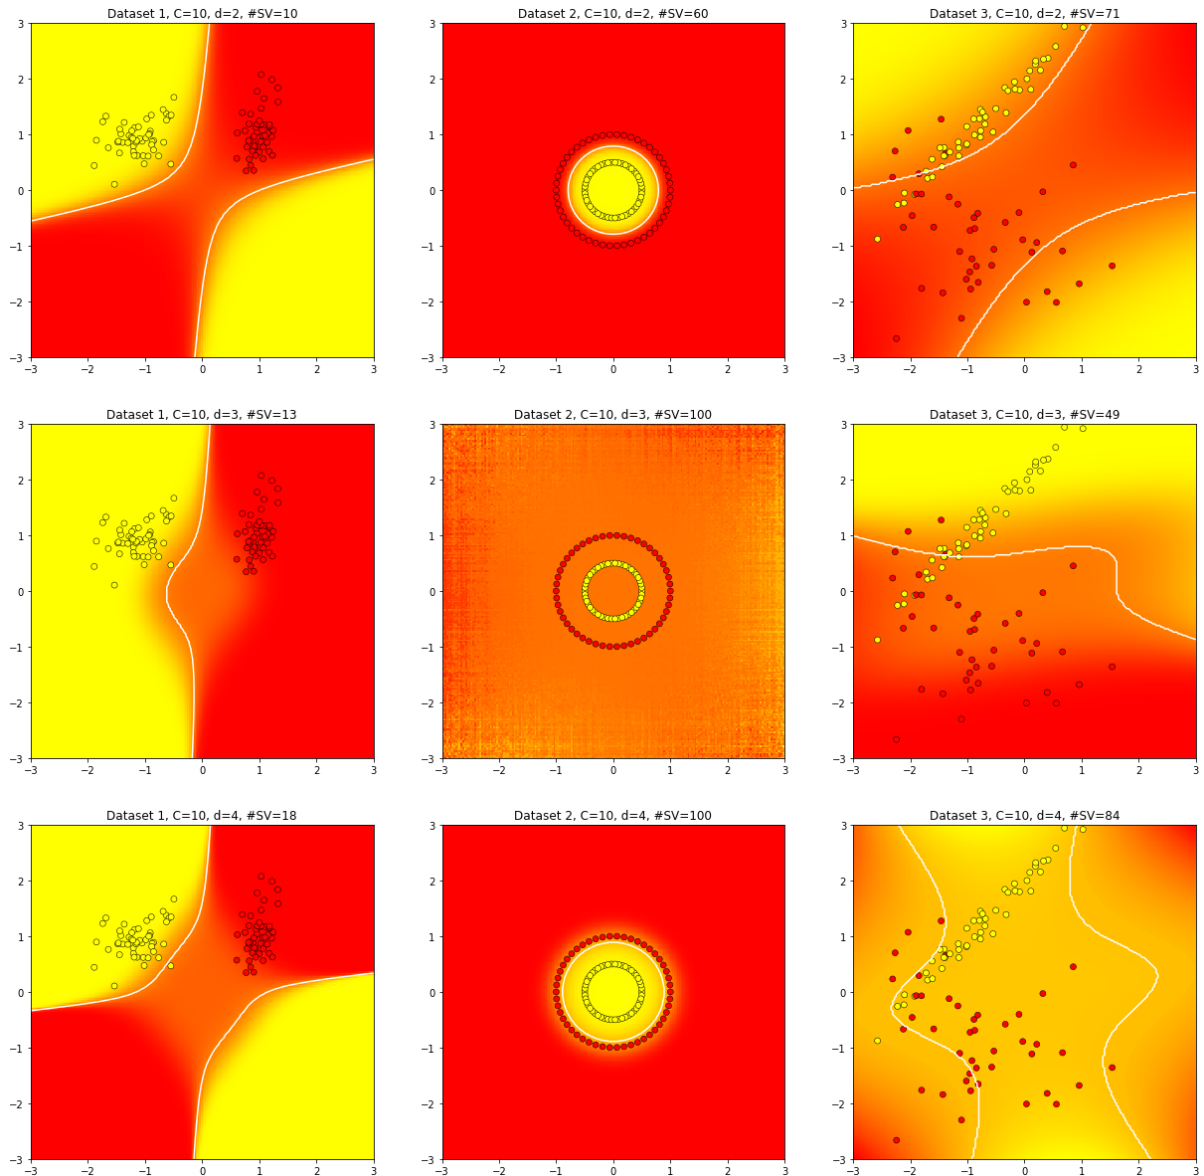
        model = SVC(kernel='poly', degree=d, gamma='auto', probability=True)
        #-----

        model.fit(X[i], y[i])
        plt.scatter(X[i][:, 0], X[i][:, 1], c=y[i], cmap='autumn', linewidths=0.5, edgecolors=(0, 0, 0, 1))
        xvals = np.linspace(-3, 3, 200)
        yvals = np.linspace(-3, 3, 200)
        xx, yy = np.meshgrid(xvals, yvals)
        zz = np.reshape(model.predict_proba(np.c_[xx.ravel(), yy.ravel()])[:, 1] * 2 - 1, xx.shape)
        plt.pcolormesh(xx, yy, zz, cmap='autumn', zorder=0)
        plt.contour(xx, yy, zz, levels=(-1., 0., 1.), colors='w', linewidths=1.5, zorder=1, linestyle='solid')

        plt.xlim([-3, 3])
        plt.ylim([-3, 3])
        plt.title('Dataset {}, C={}, d={}, #SV={}'.format(i + 1, C, d, len(model.support_vectors_)))

```

C:\miniconda\envs\cvx\lib\site-packages\matplotlib\contour.py:1243: UserWarning: No contour levels were found within the data range.
 warnings.warn("No contour levels were found")



4 Written Problems

Problem 4.1 Dual Representations [10 pts]

Read section 6.1 in Bishop, and work through all of the steps of the derivations in equations 6.2-6.9. You should understand how the derivation works in detail. Write down your understanding.

Problem 4.2 Kernels [10 pts]

Read Section 6.2 and Verify the results (6.13) and (6.14) for constructing valid kernels.

Problem 4.3 Maximum Margin Classifiers [10 pts]

Read section 7.1 and show that, if the 1 on the right hand side of the constraint (7.5) is replaced by some arbitrary constant $\gamma > 0$, the solution for maximum margin hyperplane is unchanged.

$$t_n(w^T \phi(x_n) + b) \geq 1$$

Problem 4.1

$$J(w) = \frac{1}{2} \sum_{n=1}^N \{w^T \phi(x_n) - t_n\}^2 + \frac{\lambda}{2} w^T w \quad (6.2)$$

$$\nabla_w J(w) = \frac{1}{2} \sum_{n=1}^N 2 \{w^T \phi(x_n) - t_n\} \cdot \phi(x_n) + \lambda w = 0$$

$$\begin{aligned} \Rightarrow w &= -\frac{1}{\lambda} \sum_{n=1}^N \{w^T \phi(x_n) - t_n\} \cdot \phi(x_n) = \\ &= \sum_{n=1}^N \left(\frac{\{w^T \phi(x_n) - t_n\}}{-\lambda} \right) \cdot \phi(x_n) = \sum_{n=1}^N a_n \phi(x_n) = \\ &= \Phi^T a \quad (6.3) \quad a = (a_1, a_2, \dots, a_N)^T \end{aligned}$$

we define $a_n = \left(-\frac{1}{\lambda} (w^T \phi(x_n) - t_n) \right)$ (6.4)

now we substitute $w = \Phi^T a$ into $J(w)$ to obtain

$$\begin{aligned} J(\Phi^T a) &= \frac{1}{2} \sum_{n=1}^N \left[(\Phi^T a)^T \cdot \phi(x_n) - t_n \right]^2 + \frac{\lambda}{2} \cdot (\Phi^T a)^T (\Phi^T a) \\ &= \frac{1}{2} \sum_{n=1}^N \left[a^T \Phi \cdot \phi(x_n) - t_n \right]^2 + \frac{\lambda}{2} \cdot a^T \Phi \Phi^T a \\ &= \frac{1}{2} \sum_{n=1}^N \left[a^T \Phi \cdot \phi(x_n) - t_n \right]^T \left[a^T \Phi \cdot \phi(x_n) - t_n \right] + \frac{\lambda}{2} a^T \Phi \Phi^T a \\ &= \frac{1}{2} a^T \Phi \Phi^T \Phi \Phi^T a - a^T \Phi \Phi^T t + \frac{1}{2} t^T t + \frac{\lambda}{2} a^T \Phi \Phi^T a \quad (6.5) \end{aligned}$$

we now define the Gram matrix $\underline{\underline{K}} = \underline{\underline{\Phi}} \underline{\underline{\Phi}}^T$

where the (n,m) element is

$$K_{n,m} = \phi(x_n)^T \phi(x_m) = \kappa(x_n, x_m) \quad \underline{(6.6)}$$

\Rightarrow we can rewrite $J(a)$ as

$$\begin{aligned} J(a) &= \frac{1}{2} a^T (\underline{\underline{\Phi}} \underline{\underline{\Phi}}^T) (\underline{\underline{\Phi}} \underline{\underline{\Phi}}^T) a - a^T (\underline{\underline{\Phi}} \underline{\underline{\Phi}}^T) t + \frac{1}{2} t^T t + \frac{\lambda}{2} a^T (\underline{\underline{\Phi}} \underline{\underline{\Phi}}^T) a \\ &= \frac{1}{2} a^T K K a - a^T K t + \frac{1}{2} t^T t + \frac{\lambda}{2} a^T K a \quad \underline{(6.7)} \end{aligned}$$

we now compute $\nabla_a J(a)$ and set it equal to 0

$$\Rightarrow \nabla_a J(a) = K K a - K t + \lambda K a = 0$$

$$\Rightarrow K K a + \lambda K a = K t$$

$$K a + \lambda a = t$$

$$(K + \lambda I_n) a = t$$

$$a = (K + \lambda I_n)^{-1} t \quad \underline{(6.8)}$$

We now substitute (6.8) into the linear Regression model given by $y(x) = w^T \phi(x)$

$$\text{true } w = \bar{\Phi}^T a$$

$$\Rightarrow y(x) = w^T \phi(x) = a^T \bar{\Phi} \phi(x) = \kappa(x)^T (\kappa + \lambda I_N)^{-1} t$$

(6.9)

Problem 4.2

a) consider the definition of a kernel given by $\kappa(x, x') = \phi(x)^T \phi(x')$, where $\phi(x)$ - are a non-linear feature space mapping.

consider $\kappa_1(x, x') = \phi_1(x)^T \phi_1(x')$, a kernel

now consider $\kappa(x, x') = c \cdot \kappa_1(x, x')$, if κ is a valid kernel we must be able to find a valid representation using some feature vectors

if we take $\phi(x) = c^{\frac{1}{2}} \phi_1(x)$ we can rewrite

$$\begin{aligned}\kappa(x, x') &= c^{\frac{1}{2}} \phi_1(x)^T \cdot c^{\frac{1}{2}} \phi_1(x') = c \cdot \phi_1(x)^T \phi_1(x') \\ &= c \cdot \kappa_1(x, x')\end{aligned}$$

$\Rightarrow c \cdot \kappa_1(x, x')$ - is a valid kernel

b) Now show that $\kappa(x, x') = f(x) \kappa_1(x, x') f(x')$ is a valid kernel, where $f(x) \rightarrow$ any function we can rewrite $\kappa(x, x')$

$$\kappa(x, x') = f(x) \cdot \phi_1(x)^T \phi_1(x') \cdot f(x')$$

if we define $v(x) = \phi_1(x) \cdot f(x)$, we can write

$$\kappa(x, x') = v(x)^T v(x')$$

$\Rightarrow \kappa(x, x')$ can be written as an scalar product of feature mappings \Rightarrow

$\kappa(x, x') = f(x) \kappa_1(x, x') f(x')$ is a valid kernel

Problem 9.3

suppose $\exists w_1, b_1$ that satisfy the solution to equation (7.5) and minimize equation (7.3).

Our constraint will have the form

$$t_n (w_1^T \phi(x_n) + b_1) \geq 1$$

then suppose we have a constraint

$$t_n (w^T \phi(x_n) + b) \geq \gamma, \quad \gamma \geq 0$$

if we rescale our variables such that

$w = \gamma w_1$ and $b = \gamma b_1$, we will still minimize (7.3) and our constraint would become

$$t_n ((\gamma w_1)^T \phi(x_n) + \gamma b_1) \geq \gamma$$

$$\cancel{\gamma} \cdot t_n (w_1^T \phi(x_n) + b_1) \geq \cancel{\gamma}$$

$$t_n (w_1^T \phi(x_n) + b_1) \geq 1$$

therefore replacing 1 on the r.h.s with $\gamma > 0$, does not change the solution for the maximum margin hyperplane